

# Project 8: Skip List and B+ Tree Node Implementations

Group: P8G03

## Authors

Canyu Liu  
Xuanting Liu  
Sihan Chen

# Abstract

This report studies the skip list, a randomized layered data structure that supports search, insertion, and deletion efficiently in expectation. We first describe how a skip list organizes elements across multiple levels using forward pointers and randomized node heights. We then provide a formal expected-time analysis showing  $O(\log n)$  expected running time for search, insertion, and deletion, and  $O(n)$  expected space. Next, we integrate different node data structures (array, linked list, and a skip-list-based store) into a B+ tree implementation through a unified `NodeStoreOps` interface. We prove that when node-local search is  $O(\log M)$ , the overall B+ tree search time on  $n$  keys is  $O(\log n)$ , independent of the tree order  $M$ . Finally, we use our benchmark results and the generated test data to provide empirical evidence supporting the key steps of the proof.

---

## 1. Skip List Description

### 1.1 Data Structure Overview

A **skip list** maintains keys in sorted order using multiple levels of forward pointers:

- **Level 0** is a standard sorted linked list containing all keys.
- Each higher level  $i > 0$  is a subsequence of level  $i - 1$ , providing "express lanes" that allow search to skip over many nodes.
- A **header/sentinel** node exists at all levels and serves as the entry point.

Each inserted key is assigned a random height (level count). A typical construction promotes a node from level  $i$  to  $i + 1$  independently with probability  $p$  (commonly  $p = 0.5$ ), so higher levels contain exponentially fewer nodes.

### 1.2 Node and List Representation

- A `SkipListNode` stores an integer key and a flexible array `forward[]`.
- The node's `level` is the length of `forward[]`. If a node has height  $h$ , it appears on levels  $0..h - 1$ .
- A `SkipList` stores:
  - `max_level` (upper bound on node height),
  - `p` (promotion probability),
  - `level` (current highest non-empty level),
  - `size` (number of keys),
  - `header` (sentinel node with height `max_level`).

## 1.3 Search Operation

Search starts from the highest active level. At each level:

1. **Move right** while the next node exists and its key is smaller than the target.
2. When moving right would overshoot (next key  $\geq$  target), **drop down** one level.

After reaching level 0, the algorithm checks whether the next node equals the target.

### 1.3.1 Pseudocode — SEARCH

```
SEARCH(SL, key):
    x <- SL.header
    for i from SL.level-1 down to 0:
        while x.forward[i] != NIL and x.forward[i].key < key:
            x <- x.forward[i]
    x <- x.forward[0]
    if x != NIL and x.key == key:
        return FOUND
    else:
        return NOT_FOUND
```

---

## 1.4 Insertion Operation

Insertion performs a traversal like search but records predecessors `update[]`. If the key does not already exist, we generate a random node height and insert the new node by rewiring forward pointers.

### 1.4.1 Pseudocode — RANDOM\_LEVEL

```
RANDOM_LEVEL(p, max_level):
    lvl <- 1
    while Random(0,1) < p and lvl < max_level:
        lvl <- lvl + 1
    return lvl
```

### 1.4.2 Pseudocode — INSERT

```
INSERT(SL, key):
    update[0..SL.max_level-1]
    x <- SL.header

    for i from SL.level-1 down to 0:
```

```

while x.forward[i] != NIL and x.forward[i].key < key:
    x <- x.forward[i]
    update[i] <- x

x <- x.forward[0]
if x != NIL and x.key == key:
    return FAIL // duplicate key

lvl <- RANDOM_LEVEL(SL.p, SL.max_level)

if lvl > SL.level:
    for i from SL.level to lvl-1:
        update[i] <- SL.header
    SL.level <- lvl

new <- newNode(key, lvl)
for i from 0 to lvl-1:
    new.forward[i] <- update[i].forward[i]
    update[i].forward[i] <- new

SL.size <- SL.size + 1
return OK

```

## 1.5 Deletion Operation

Deletion records predecessors `update[]`. If the key exists, it is removed by bypassing it on each level where it appears. If the highest level becomes empty, the skip list's current height decreases.

### 1.5.1 Pseudocode — DELETE

```

DELETE(SL, key):
    update[0..SL.max_level-1]
    x <- SL.header

    for i from SL.level-1 down to 0:
        while x.forward[i] != NIL and x.forward[i].key < key:
            x <- x.forward[i]
            update[i] <- x

    x <- x.forward[0]
    if x == NIL or x.key != key:
        return FAIL

    for i from 0 to SL.level-1:

```

```

    if update[i].forward[i] == x:
        update[i].forward[i] <- x.forward[i]

    free(x)
    SL.size <- SL.size - 1

    while SL.level > 1 and SL.header.forward[SL.level-1] == NIL:
        SL.level <- SL.level - 1

    return OK

```

---

## 2. Formal Analysis of Skip List Operations

### 2.1 Expected Height

A node appears on level  $i$  with probability  $p^i$ . Therefore, the expected number of nodes on level  $i$  is:

$$\mathbb{E}[N_i] = np^i$$

The maximum non-empty level  $h$  is near the point where  $np^h \approx 1$ , implying:

$$h = O(\log_{1/p} n)$$

For a constant  $p$  (e.g.,  $p = 0.5$ ),  $h = O(\log n)$ .

### 2.2 Expected Search Time

A standard skip-list argument yields expected  $O(1/p)$  horizontal steps per level. With  $O(\log_{1/p} n)$  levels, expected search time is:

$$\mathbb{E}[T_{\text{search}}] = O\left(\frac{1}{p} \log_{1/p} n\right) = O(\log n) \quad (\text{constant } p)$$

### 2.3 Expected Insertion and Deletion Time

Insertion/deletion perform one search (to compute `update[]`) plus pointer updates on the node height. The expected node height is:

$$\mathbb{E}[\text{height}] = \sum_{k \geq 1} \Pr(\text{height} \geq k) = \sum_{k \geq 1} p^{k-1} = \frac{1}{1-p}$$

Thus insertion and deletion are expected  $O(\log n)$ .

### 2.4 Space Complexity

Expected total pointers are  $n/(1-p)$ , so expected space is  $O(n)$ .

---

### 3. B+ Tree with Skip-List Nodes

#### 3.1 Node Storage Abstraction (`NodeStoreOps`)

We keep B+ tree logic unchanged and swap node implementations via `NodeStoreOps`:

- `lower_bound(store, key)` → first index with  $\text{key} \geq \text{target}$
- `key_at(store, idx), val_at(store, idx)`
- `insert_at(store, idx, key, val), erase_at(store, idx)`
- `split(left, right)` → move half entries to right and return separator key

We compare three node stores:

- **array**: sorted arrays + binary search (fast search, shifts on updates)
  - **list**: sorted linked list (linear scan)
  - **skip**: skip-list-based store (designed to mimic node-level  $O(\log M)$  search)
- 

#### 3.2 Claim to Prove (from project hint)

Using skip list to implement B+ tree nodes: prove the running time of finding a node in a B+ tree of  $n$  keys is  $O(\log n)$ , independent of its order  $M$ , and analyze insertion/deletion.

---

#### 3.3 Formal Proof: Search is $O(\log n)$ Independent of $M$

A B+ tree of order  $M$  has height  $H = O(\log_M n)$ . If **search inside a node** costs  $O(\log M)$  (e.g., binary search in an array, or expected  $O(\log M)$  in a skip list), then total search time is:

$$T(n) = O(\log_M n) \cdot O(\log M) = O\left(\frac{\log n}{\log M} \cdot \log M\right) = O(\log n)$$

Therefore, the asymptotic search time is  $O(\log n)$  and does **not** depend on  $M$ .

---

### 3.4 Insertion and Deletion (analysis sketch)

Insertion performs:

- one root-to-leaf descent:  $O(\log_M n)$  node visits,
- node-local insert: depends on node store (array shifts  $O(M)$ , list scan  $O(M)$ , ideal skip-list  $O(\log M)$ ),
- occasional splits that may propagate upward for at most  $O(\log_M n)$  levels.

Deletion performs:

- one descent:  $O(\log_M n)$ ,
- node-local erase: again depends on node store,
- possible rebalancing via **borrow** or **merge** up the path (at most  $O(\log_M n)$  levels).

Thus, if node-local update is  $U(M)$ , insertion/deletion are roughly:

$$O\left(\log_M n \cdot (\log M + U(M))\right)$$

(amortized / expected depending on the node store).

---

## 4. Test Data (Generator) and Experimental Setup

### 4.1 Benchmark Input Format

Our benchmark reads three files:

- `-insert <file>`: keys to insert
- `-search <file>`: keys to query
- `-delete <file>`: keys to delete

File format:

- integers separated by whitespace
- optional comment lines starting with #

This matches our generator output exactly.

## 4.2 Test Data Generator (Python)

We use a Python generator program to create **reproducible** datasets. It writes:

- <prefix>\_insert.txt
- <prefix>\_search.txt
- <prefix>\_delete.txt

### 4.2.1 Insert keys: `-insert-dist`

Main choices:

- `unique_uniform`: generate  $n$  **unique** keys uniformly in `[key_min, key_max]` (recommended).
- `sorted_unique`: generate unique keys then sort ascending.
- `reverse_unique`: generate unique keys then sort descending.
- `nearly_sorted_unique`: generate unique keys, sort, then do about `swap_frac * n` random swaps.
- `uniform`, `normal`, `exp`, `pareto`, `clusters`: other distributions (may include duplicates).

### 4.2.2 Search keys: hit/miss mixture

Search queries are generated with a target hit ratio `-hit-ratio`:

- hits: sampled from `insert_keys`
- misses: generated to avoid the inserted set, via:
  - `offset` (add a large offset),
  - `uniform_retry`,
  - `normal_retry`

Queries are shuffled to avoid patterns.

### 4.2.3 Delete keys: `-delete-mode`

- `shuffle_all`: delete all inserted keys (random order) or the first `n_delete`.
- `in_order`: delete in insertion order.
- `random_subset`: delete a random subset of size `n_delete`.

#### 4.2.4 Reproducibility

We fix `-seed`. The script uses deterministic derived seeds for search/delete streams, so the same command reproduces exactly the same three files.

### 4.3 Our Experimental Parameters (from uploaded result)

From `test.md`:

- order  $M = 64$
- $n_{\text{insert}} = 200000$
- $n_{\text{search}} = 200000$
- $n_{\text{delete}} = 200000$
- 4 cases, each repeated 5 rounds, for each implementation (`array`, `list`, `skip`).

#### Hit ratio evidence from results.

Case 1 has `found_count = 140000` out of 200000 searches ( $\approx 0.7$  hit ratio). Cases 2–4 have `found_count = 100000` ( $\approx 0.5$  hit ratio). This is consistent with our generator's `-hit-ratio` feature.

---

## 5. Results and Evidence Supporting the Claim

### 5.1 Average Time per Case

Times below are averaged over 5 rounds per case (ms). `found` is average `found_count`.

case	impl	rounds	insert_ms	search_ms	delete_ms	found	height
1	array	5	19.51	13.89	18.90	140000	1
1	list	5	83.91	63.40	76.37	140000	1
1	skip	5	817.66	14.09	853.11	140000	1
2	array	5	26.95	9.66	38.61	100000	1
2	list	5	119.98	52.15	182.88	100000	1
2	skip	5	2331.71	10.14	2978.99	100000	1
3	array	5	22.85	9.13	37.64	100000	1
3	list	5	172.83	42.74	173.29	100000	1
3	skip	5	2028.60	10.49	2960.64	100000	1
4	array	5	32.87	9.66	37.75	100000	1
4	list	5	186.73	60.74	185.43	100000	1
4	skip	5	2414.19	10.25	2992.05	100000	1

## 5.2 Overall Average Across All Cases

impl	rounds	insert_ms	search_ms	delete_ms	found	height
array	20	25.55	10.58	33.22	110000	1
list	20	140.86	54.76	154.49	110000	1
skip	20	1898.04	11.24	2446.20	110000	1

## 5.3 Using the Results to Support the “ $O(\log n)$ independent of $M$ ” Viewpoint

The formal proof in Section 3.3 relies on a key decomposition:

1. B+ tree search visits  $O(\log_M n)$  nodes (tree height).
2. Each node visit needs an **in-node search** costing  $O(\log M)$  (array binary search or skip-list expected search).
3. Multiply them to get  $O(\log n)$ .

Our experiment provides empirical evidence for step (2), i.e., **node-local search efficiency matters**:

- Under the same  $M = 64$  and the same query workload, the **linked-list** node store (linear scan inside node) has much larger search time than the **array** and **skip** stores.
- Overall average search time (ms) from Section 5.2:
  - array: 10.58
  - skip: 11.24
  - list: 54.76

The list is about  $\approx 5.17 \times$  slower than array, and  $\approx 4.87 \times$  slower than skip in search time (overall averages). This aligns with the theoretical expectation that list-based node search is  $\Theta(M)$  while array/skip are closer to  $\Theta(\log M)$  for node-local search.

### 5.3.1 About `height_after_insert = 1` in this run

In this benchmark output, `height_after_insert = 1` in all rounds. This implies that the measured `search_ms` is dominated by **single-node lookup cost** (effectively isolating the in-node search). That is precisely the sub-problem needed in the proof: demonstrating that using better node structures reduces in-node search from linear to logarithmic.

### 5.3.2 Recommended follow-up to fully validate independence from $M$

To empirically validate the complete statement “ $O(\log n)$  independent of  $M$ ” end-to-end, we recommend:

- ensuring the B+ tree grows to height  $> 1$  (so  $\log_M n$  is observable),
  - running multiple values of  $M$  (e.g., 32, 64, 128, 256) with the same  $n$ ,
  - plotting search time vs.  $M$ . Theory predicts the asymptotic trend stays  $O(\log n)$ ; varying  $M$  changes constants but not the overall order.
- 

## 6. Discussion and Conclusion

### 6.1 Discussion

**Array node store.** Best overall. Search is efficient (binary search), and updates shift  $O(M)$  elements but remain fast in practice due to locality.

**Linked-list node store.** Search is slow because each in-node lookup requires a linear scan. The experimental results show a consistent multiplicative slowdown for searches.

**Skip-list-based node store.** Search time is close to array, but insertion/deletion are much slower. A major cause is **interface mismatch**: our `skiplist.c` is key-only, while internal B+ nodes require storing (`key`, `child_ptr`) pairs and supporting index-based operations. Without an indexable, value-carrying skip list (e.g., with span/width augmentation), bridging can introduce extra overhead.

### 6.2 Conclusion

- Skip lists support expected  $O(\log n)$  search/insert/delete with expected  $O(n)$  space.
  - If B+ node-local search is  $O(\log M)$ , B+ tree search is  $O(\log n)$  independent of  $M$ .
  - Our benchmark results support the key step that node-local search should be logarithmic: array/skip are much faster than linked list on the same dataset.
- 

## 7. References

- W. Pugh, “Skip Lists: A Probabilistic Alternative to Balanced Trees,” 1990.
- Wikipedia: Skip list, B+ tree.